

Model Context Protocol (MCP)

A standardized way for AI applications to connect with tools and data

What is MCP?

Model Context Protocol is an open protocol designed to standardize how AI applications interact with external tools, resources, and data sources. Instead of building a separate integration pattern for every model and application, MCP provides a common structure through which an AI system can discover and use capabilities exposed by an MCP server.

Client and server model

An MCP-enabled application acts as the host or client side of the interaction, while an MCP server exposes capabilities. A server can provide tools that perform actions, resources that expose information, and prompts or other structured capabilities. The application decides when those capabilities are appropriate to use.

Tools

Tools represent callable operations. For example, an MCP server could expose database operations, file search, API access, or a calculation function. A tool normally has a defined name, description, and input schema so the consuming application can understand how to call it safely and consistently.

Resources

Resources provide contextual information to an AI application. They can represent documents, database information, files, or other readable content. Keeping data access separate from the model itself makes integrations more modular and allows the same capability to be reused by different AI applications.

Why MCP is useful

MCP reduces the need for custom, tightly coupled integrations between AI applications and every external system. It encourages a clear boundary between reasoning and external capabilities, making tool discovery and integration more standardized. Security, permissions, validation, and careful tool design remain essential because an exposed tool can perform real operations.

Reference document for the Veritas portfolio demonstration.